




How to Tune Your Models








It's All Downhill From Here




This image was generated using AI tools





This image was generated using AI tools




Loss Functions

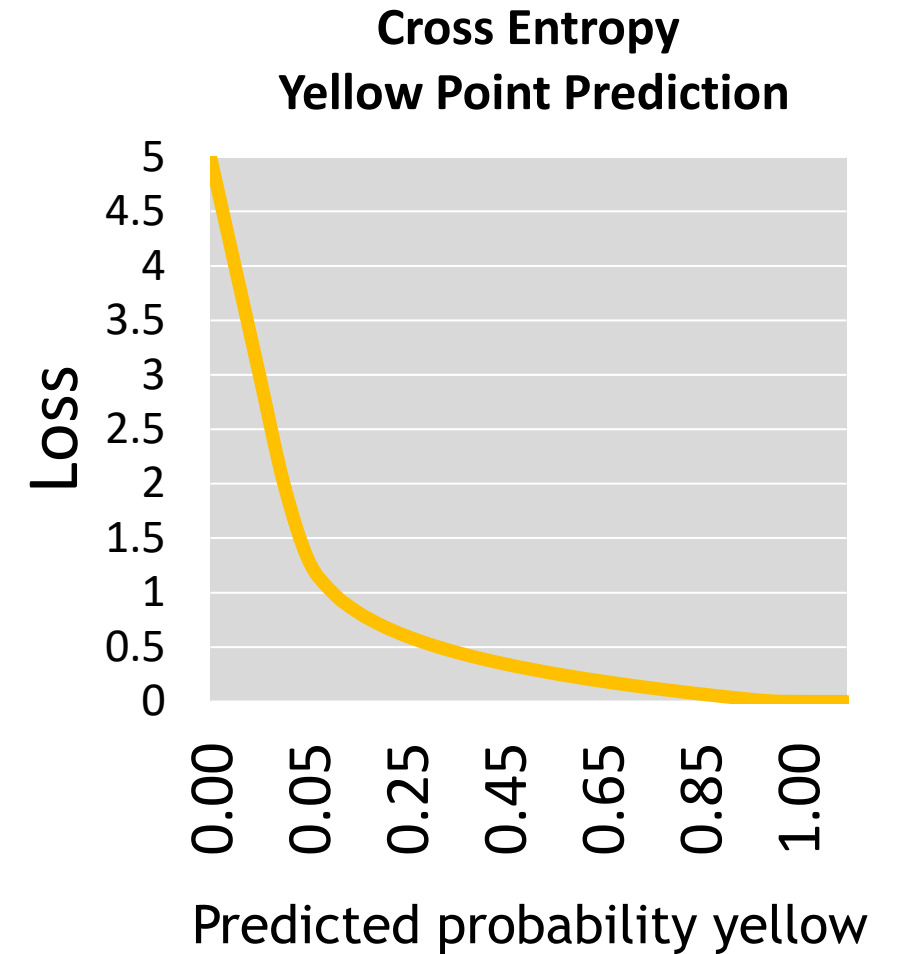

How do we quantify
prediction error?




Mean Squared Error (MSE)




Cross Entropy








Training a neural network with Stochastic Gradient Descent.

SGD, Batch GD, and Mini-Batch GD on a Complex Cost Function




This image was generated using AI tools

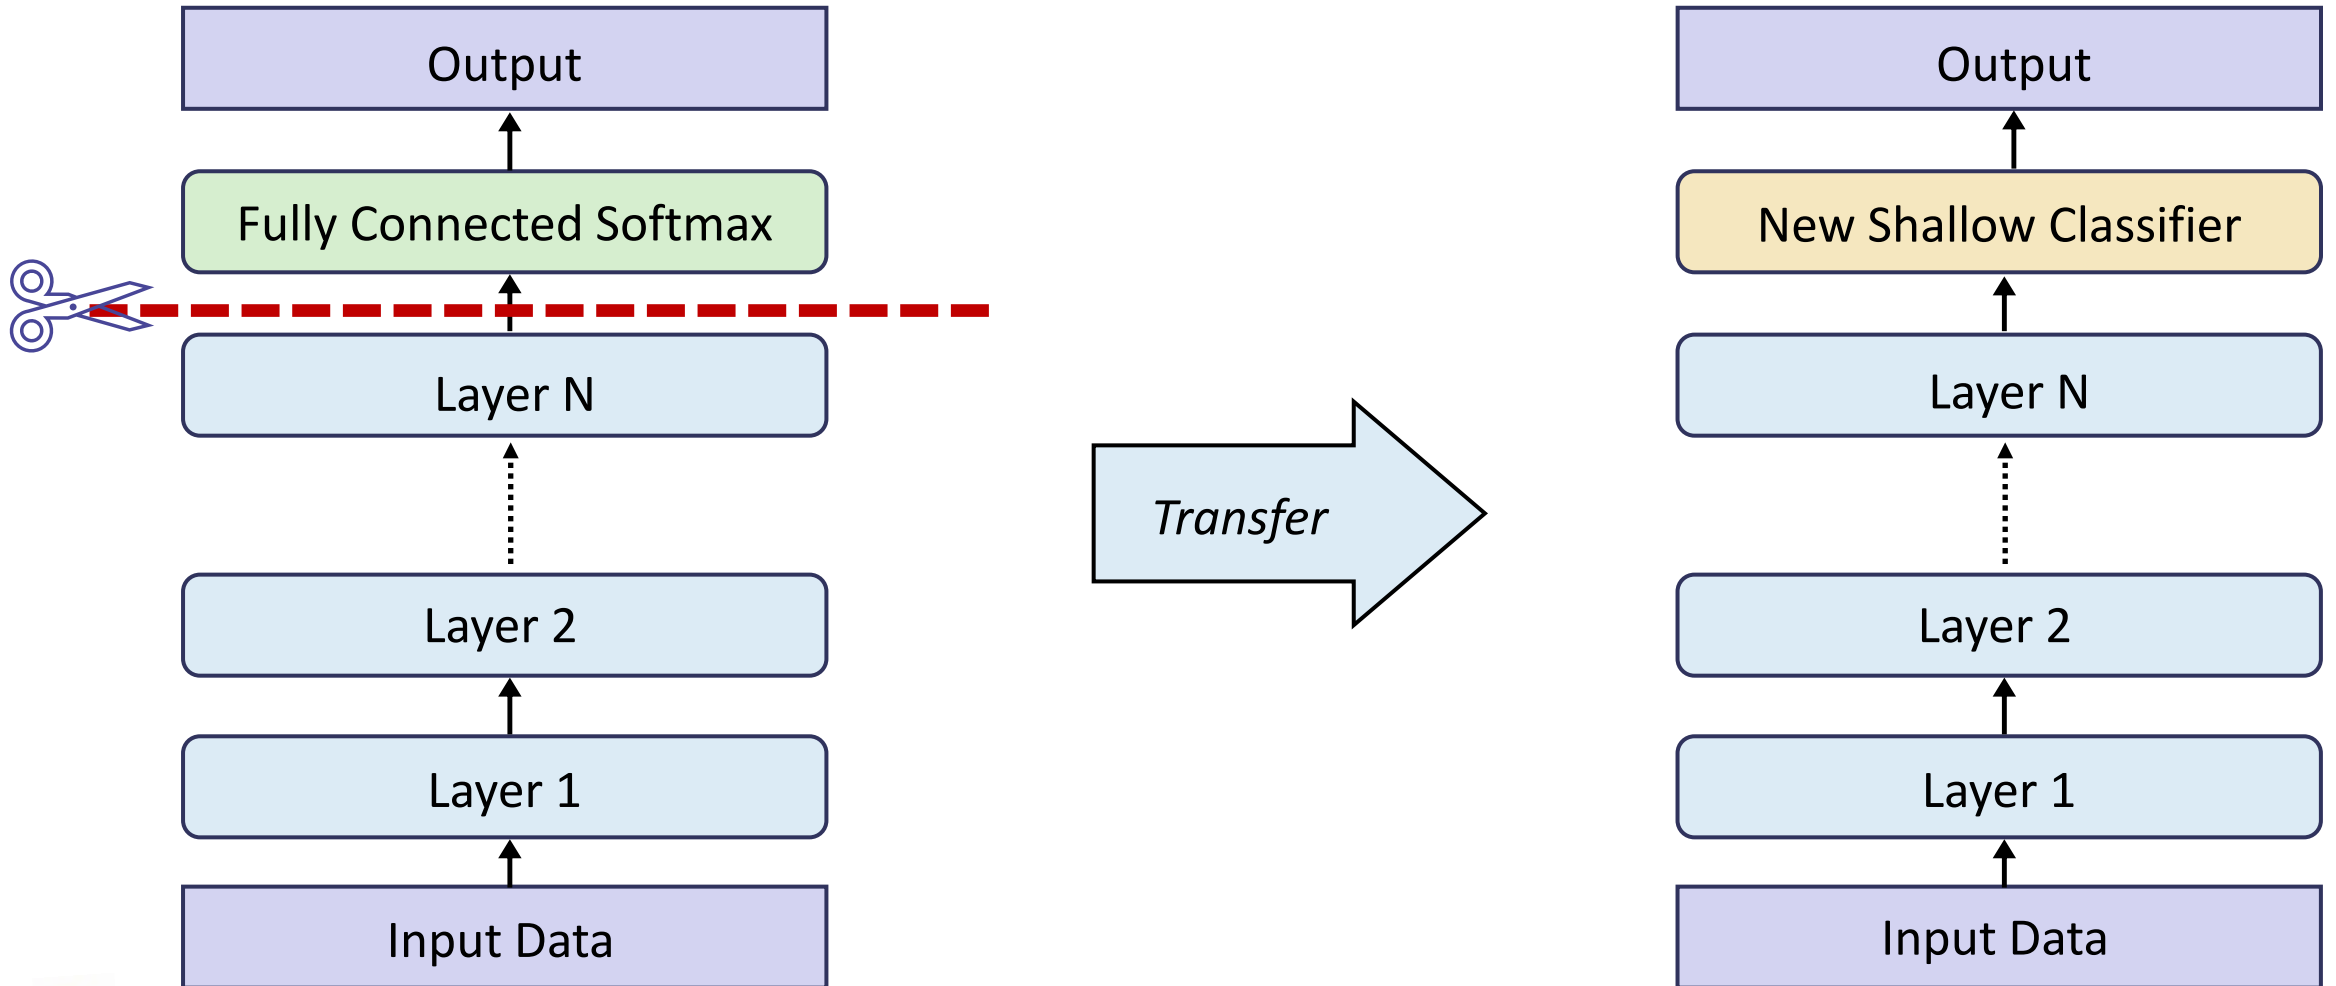




This Image Was Generated Using AI Tools



We Can Also Use What We've "Learned"





Hyperparameter Optimization

03_bees_vs_wasps.ipynb

This notebook will walk you through building and training your own image classification model, then allow you to compare different hyperparameter optimization configurations!

